

The *Rabbit Hole*

Two essays on trust architecture, self-selection, and the assumption worth replacing.

I. The Rabbit Hole Was the Audit

Rick Astley didn't plan his comeback. He was a one hit wonder who became a stay at home dad and occasionally worked as a waiter. Someone embedded one of his old music videos in a hyperlink, called it something else, and the internet played along for years. Somewhere in there he became genuinely beloved again. Not because he managed his reputation carefully, but because the joke was good and he was in on it. Glastonbury stage. Thirty years later. A second chance nobody saw coming, least of all Rick.

Never gonna give you up.

(I once Rick Rolled two division presidents named Rick in a mass email. Nobody got it. Still worth it.)

I told everyone it existed.

I posted the invitation on Substack — the egg, the Anne Murray reference, the instructions. I said most people would never find it. I said you just have to decide if you are that person. And then I watched.

The Strategic Linguist found it and named what it was doing without being briefed. Dr. Sam Illingworth discovered his own work embedded inside it and thanked me publicly. A handful of others followed all the way down. Most people, even with the map in hand, didn't go.

That last part is the interesting data point.

When I launched my website, I made myself one rule: almost everything is evergreen. The writing updates. The podcast updates. Everything else gets reviewed once a year and otherwise left alone. But I wanted one layer that was alive. That did something. So I hid an egg. (Bonus points if you hum Anne Murray's Teddy Bear Picnic while hunting.)

Find it and you're invited to follow the rabbit. Six depths of ideas about trust, about signals, about what it means to notice something real when everything around you has been optimized to look real but

isn't. Each page has a depth counter: DEPTH X OF 6. And a button: Keep going.

The counter is the rope. You always know where you are. The button is always a choice. Nobody falls accidentally. You have to keep deciding to go further.

At the bottom: three tools I genuinely believe in. A game for spotting AI-hallucinated citations (it will humble you). A critical thinking bot, because the skill needs practice, not just awareness. A Substack on how language shapes power, because if you're reading trust signals, you need a linguist.

By the time you get there, you've done the thing. Not read about trust signal literacy. Practiced it. I didn't lecture you about how I think. You just thought the way I think, for six pages. That was the design: make you do the thing before you realize you're being taught anything.

The rabbit hole isn't permanent. It's seasonal. It went up for Easter and it will come down. Then another egg, another window, different tools and different people. Each iteration is a deliberate choice: whose work belongs at depth 6 right now? Who should know they're in it?

Dr. Illingworth didn't receive a collaboration pitch. He received a public demonstration that someone took his work seriously enough to put it at the bottom of six depths of ideas about trust. That's a different professional signal entirely.

Here's something I've been working out: trust moves differently across in-person, digital, and AI registers. The Easter window was the in-person layer. Real scarcity. Real time. You had to be there. The Substack thread, the screenshots, this essay — that's digital. The artifact of the experience outlasting the experience itself. A bot scanning for her IP after the window closes is AI. It finds the record of something it couldn't have been present for. The bot detects. She witnesses. Those are not the same thing. And that distinction is exactly what the rabbit hole is about.

I spend a lot of time helping organizations surface the assumptions operating beneath their decisions. The ones nobody examines because everybody shares them. I'm apparently not immune to my own methodology.

The assumption I was carrying: that explaining my thinking was how people would come to understand it. Clearer articulation. Better case studies. More explicit frameworks.

The rabbit hole tests a different hypothesis. You don't need to explain how you think if you can build something that is the thinking, then let

people find out for themselves whether they can follow it.

The people who go all the way down are exactly who I want in the room. Not because they passed a test. Because they demonstrated something about how they pay attention that no credentials or categories could have established.

No cold outreach. No pitch. Embed someone's work in a trust architecture because you genuinely believe it belongs there. Make it seasonal. Let them discover it. The reputation compounds through precision and genuine inclusion, not visibility and volume.

The egg goes back in the ground. The website goes quiet. Somewhere between now and next time, I'll figure out what belongs at depth 6 for the next window.

That's the practice. Not a prediction.

This essay is part of the Witnessed Trust series. The Assumption-Ground Audit is a methodology for surfacing expiring assumptions before decisions harden into policy.

II. The Rabbit Hole Method

The mistake most people make when building a reputation is designing for visibility. I made it for years.

More content, more consistency, more presence. Optimize for reach and call it strategy. What the approach rests on, underneath — the assumption nobody examines because everybody shares it — is that explaining yourself is how understanding transmits. That clearer articulation is the answer. Better case studies, more organized arguments, more frequent proof.

That assumption is expiring. The Rabbit Hole Method starts by naming it.

Before you build anything, you audit the premise you're replacing. Not "I want to try something different." That's a mood. The specific operating assumption beneath your current approach — stated precisely enough that you'd recognize it if someone else said it back to you. Once you can say it out loud, you can test something else.

Then you commit the architecture. Almost everything is evergreen. The writing updates, the podcast updates, everything else gets reviewed once a year and left alone. Inside that committed structure, one or two designated play spaces — alive, iterative, deliberate. The

constraint is the design. Without it you iterate constantly and never find out what holds. The playground is only useful if you also have the monument it sits inside.

Then you design the descent.

Six depths. What lives at each level? What's the surface provocation — the thing that makes someone curious enough to take the first step? Where does the thesis sit? Not at the top. Make them approach it. Every page needs a rope: a depth counter, an orientation marker, a button that makes the choice to continue explicit. DEPTH X OF 6. Keep going. Nobody falls accidentally. You have to keep deciding to go further. That's not a UX feature. It's a structural argument about what trust requires.

The bottom is where the method's integrity lives.

Not what's popular. Not what signals your affiliations or improves your associations. Whose work genuinely belongs in this argument, at this depth, right now? The people who reach the bottom have demonstrated how they pay attention. What they find there has to be worthy of that demonstration. Get this right and people feel found. Get it wrong and they feel misled. Those are very different experiences and you only discover which one you built after the window closes.

Then you set the window.

Real scarcity changes the quality of attention something receives. The rabbit hole goes up for Easter. It comes down. A seasonal window is also an iteration cycle — each time it returns, you update the tools you trust, whose work belongs at depth 6, which ideas you're currently testing. Different people, different resources, different window. The method stays consistent. The content evolves because you evolved.

The invitation is the first filter. You tell people it exists. You don't direct them to it. "It's on my website. You just have to decide if you are that person." Most people, even with the map in hand, won't go. That's data. The ones who follow all the way down have demonstrated something that no credentials or categories could have established. The invitation was public. The architecture self-selected.

After the window closes, you archive the evidence. Screenshots. Public responses. Who followed? Who named what it was doing without being briefed? The digital layer outlasts the in-person window. The AI layer finds the record of something it couldn't have been present for. Each seasonal iteration produces feedback. The playground gets smarter because the lab returns results.

I spend a lot of time helping organizations name the assumptions operating beneath their decisions. The Assumption-Ground Audit is step one of the Rabbit Hole Method for the same reason it's the first move in every engagement: you cannot replace a premise you haven't examined. What seems like a visibility problem is almost always an assumption problem. And the fix is never more explanation.

The hypothesis the method is testing: that a trust architecture, designed carefully, self-selects the exact people you want in the room — without a single cold ask.

The assumption it's replacing: that explanation is how thinking transmits.

Stop explaining. Start building something that is the thinking. Then let people find out for themselves whether they can follow it.

This is the method behind the Witnessed Trust series. The Assumption-Ground Audit is a practice for surfacing expiring assumptions before decisions harden into policy.